

Notebook

December 29, 2019



```
In [8]: import pandas as pd
import numpy as np
import json
import matplotlib.pyplot as plt
from util import resultsmgr, analyzer

In [9]: unbalanced_results = resultsmgr.load('CensusIncome/results/1577576647.json')
balanced_results = resultsmgr.load('CensusIncome/results/1577579934.json')

def graphBestByModelType(results, **kw):
    best = analyzer.bestparams(results)
    fig, ax = plt.subplots()
    ax.set_ylim((0, 1))
    rex = ax.bar(range(len(best)), [best[b]['metrics']['testing-accuracy'] for b in best])
    ax.set_xticks(range(len(best)))
    ax.set_xticklabels(list(best))
    for r in rex:
        ht = r.get_height()
        ax.annotate(
            str(round(ht, 2)),
            xy=(r.get_x() + r.get_width() / 2, ht),
            xytext=(0, 3), # 3 points vertical offset
            textcoords="offset points",
            ha='center', va='bottom')
    ax.set_title(kw.get('title', 'Accuracies'))
    return best
```

0.1 Look at best performance for each model type

```
In [10]: unb_best = graphBestByModelType(unbalanced_results, title="Accuracies by model type\nunbalanced")
b_best = graphBestByModelType(balanced_results, title="Accuracies by model type\nbalanced")
```



0.2 It makes sense to do worse in a balanced setting, since we aren't being lazy and always guessing the predominant class. I'm going to stick with the model that performed best in the balanced setting since it is more general and "real world"

```
In [11]: print(json.dumps(b_best['LinearSVC'], indent=4))
```

```
{
  "params": null,
  "metrics": {
    "training-accuracy": 0.6142781033564199,
    "testing-accuracy": 0.8806878306878307
  }
}
```

0.3 Seems like a LinearSVC will be best so lets GridSearch on a wide range of params

```
In [13]: from util import preprocessor, splitfuncs, metrics
        from sklearn.model_selection import GridSearchCV
        from sklearn.svm import LinearSVC

        # setup model and params to test
        model = LinearSVC()
        parameters = {
            'loss': ['hinge', 'squared_hinge'],
            'C': np.linspace(0.1, 5, 5)
        }

        # get data and train/test splits
        pre = preprocessor()
        trx, try_ = pre.data
        # this is optional but lets me CV on balanced sets
        splits = []
        for i in range(5):
            tri, i, tei, j = splitfuncs.splitBalanced(trx, try_)
            splits.append((tri.index, tei.index))

        # get our best params and retrain a new model
        clf = GridSearchCV(model, parameters, cv=splits)
        gs = clf.fit(trx, try_)
        model = LinearSVC(**gs.best_params_)

        # check model performance
        trx, try_, tex, tey = splitfuncs.splitBalanced(trx, try_)
        model.fit(trx, try_)
        preds = model.predict(tex)
        acc = metrics.acc(preds, tey)
```

```
print("Best params: ", gs.best_params_)  
print("Accuracy:     ", acc)
```

```
Best params: {'C': 0.1, 'loss': 'squared_hinge'}  
Accuracy:    0.9004761904761904
```

0.4 So that didn't get us very far, but good to know we're doing the best we can